

# Fluid Simulation for Video Games (part 21)

By Dr. Michael J. Gourlay

## Recapitulation

We want games to be fun and look pretty and plausible.

Fluid simulation can augment game mechanics and enhance aesthetics and realism to video games.

Video games demand high performance on a low budget, but not great accuracy. By "low budget" I mean both computational and human resources: It has to run fast, and it can't take a lot of dev or artist time. There are many ways to simulate fluids. In this series I explained methods well-suited to video games: cheap, pretty and authorable.

If you want to simulate fluids in video games, you have many challenges to overcome. Fluid mechanics is a complicated topic soaked in mathematics, which can take a lot of time to understand. It is also numerically challenging; naïve implementations are unstable or just behave wrong. And because fluids span every point in space, simulating them costs a lot of computational resources -- both processing and memory. (While fluid simulations are well-suited to running on a GPU, in video games, the GPU tends to be completely busy with rendering.) The easiest and most obvious way to improve numerical stability (use more and smaller time steps) adds drastically more computational cost, so usually other techniques are employed, which effectively increase the viscosity of the fluid, which means that in-game fluids tend to look thick and goopy. The most popular techniques (like smoothed particle hydrodynamics) are not well-suited to delicate, wispy motions like smoke and flame. A simulation approach that met these challenges would help add more varieties of fluids (including flames and smoke) to video games.

To meet these challenges, I presented a fluid simulation technique suited to simulating fine, wispy motion, and which builds upon a particle system paradigm found in any complete 3D game engine. It can use as many CPU threads as are available -- more threads would permit more sophisticated effects.

This approach has many synergies with game development:

- Game engines support particle systems and physics simulations.
- CPU's often have unused cores.
- Fluid simulation readily parallelizes.

The approach I presented uses the Vortex Particle Method, an unusual technique which yields some benefits:

- It is a particle-based method that can reuse an existing particle engine.
- Vortex simulations are well-suited to delicate wispy motion like smoke and flame.
- The simulation algorithm is numerically stable even without viscosity, either explicit or implicit.

This series explains the math and numerics, which I recapitulate in the next section. I presented variations on the theme of particle-based fluid simulation, and included a model of thermal convection and combustion so that the system can also model flames. I showed two numerical approaches -- integral and differential -- compared their relative merits, and presented a hybrid approach that exploits the benefits of each approach

while avoiding their pitfalls. The result is a fast (linear in the number of particles) and smooth fluid simulation capable of simulating wispy fluids like flames and smoke.

Despite its apparent success in certain scenarios, the vortex particle method has some significant limitations in other scenarios. For example, it has difficulty representing interfaces between liquids and gases, such as the surface of a pool of water. I took a brief detour into Smoothed Particle Hydrodynamics to explore one way that VPM and SPH could be joined in a common hybrid framework, but (to phrase it generously) I left much room for improvement.

This series of articles leaves several avenues unexplored. I will conclude this article with a list of ideas that I encourage you to explore and share back with the community.

[Parts 1](#) and [2](#) summarized fluid dynamics and simulation techniques. [Parts 3](#) and [4](#) presented a vortex-particle fluid simulation with two-way fluid-body interactions that run in real time. [Part 5](#) profiled and optimized that simulation code. [Part 6](#) described a differential method for computing velocity from vorticity. Figure 1 shows the relationships between the various techniques and articles. [Part 7](#) showed how to integrate a fluid simulation into a typical particle system. [Parts 8, 9, 10](#) and [11](#) explained how to simulate density, buoyancy, heat and combustion in a vortex-based fluid simulation. [Part 12](#) explained how improper sampling caused unwanted jerky motion and described how to mitigate it. [Part 13](#) added convex polytopes and lift-like forces. [Parts 14, 15, 16, 17](#) and [18](#) added containers, smoothed particle hydrodynamics (SPH), liquids and fluid surfaces.

[Part 19](#) explains details on how to use a treecode algorithm to integrate vorticity to compute vector potential. Use the treecode algorithm to compute vector potential only at boundaries, as shown in Figure 2. (Later, the vector Poisson solver will “fill in” vector potentials through the domain interior.)

## Fluid Mechanics

With the practical experience of having built multiple fluid simulations behind us I now synopsise the mathematical and physical principles behind those simulations.

Fluid simulation entails running numerical algorithms to solve a system of equations simultaneously. Each equation governs a different aspect of the physical behaviors of a fluid:

### Momentum

The **momentum equation** (one version of which is the famous Navier-Stokes equation) describes how momentum evolves and how mass moves. This is a non-linear equation, and the non-linearity is what makes fluids so challenging and interesting.

Vorticity is the curl of momentum. The **vorticity equation** describes where fluid swirls -- where it moves in an "interesting" way. Since vorticity is a derivative of momentum, solving the vorticity equation is tantamount to solving the momentum equation. This series of articles exploits that connection and focuses numerical effort only where the fluid has vorticity -- which is usually much sparser than where it has momentum. Vorticity

equation also implicitly discards divergence in fluids -- the part that relates to the compressibility of fluids. Correctly dealing with compressibility requires more computational resources or more delicate numerical machinations. But in the vast majority of scenarios pertinent to video games, you can neglect compressibility. So using the vorticity equation also yields a way to circumvent the compressibility problem.

**Advection** describes how particles move. Both the momentum equation and the vorticity equation have an advective term. It is the non-linear term, so it is responsible both for the most interesting aspects of fluid motion, and also the most challenging mathematical and numerical issues. Using a particle-based method lets us separate out the advective term and handle it by simply moving particles around according to a velocity field. This makes it possible -- easy, in fact -- to incorporate the vortex particle method into a particle system. It also lets the particle system reuse the velocity field both for the "vortex particles" used by the fluid simulation and for propagating the "tracer" particles used for rendering visual effects.

The **buoyancy** term of the momentum and vorticity equations describes how gravity, pressure and density induce torque. This effect underlies how hot fluids rise, so is crucial for simulating the rising motion of flames and hot smoke. Note that the VPM simulation technique in this series did not model pressure gradients explicitly but instead assumed that pressure gradients lie entirely along the gravity direction. This let the fluid simulate buoyancy despite not having to model pressure as its own separate field. To model pressure gradients correctly typically requires modeling compressibility -- which, as mentioned elsewhere, usually costs a lot of computational resources. So by making the simplifying assumption that the pressure gradient always lies along the gravity direction permits a drastic computational savings. Computing density gradients requires knowing the relationship between adjacent particles. In this series I presented two ways to solve this: grid-based and particle-based. The grid-based approach directly employs a spatial partitioning scheme that is also used in computing viscosity effects (mentioned below). The particle-based approach uses a subset of the algorithms that Smoothed Particle Hydrodynamics uses. Both can yield satisfactory results so the decision comes down to which costs less.

The **stress & strain** terms in the momentum and vorticity equations describes how pressure and shears induce motion within a fluid. This is where viscosity enters the simulation. Varying the magnitude and form of viscosity permits the simulation to model fluids ranging from delicate, wispy stuff like smoke, or thick, goopy stuff like oil or mucous. The fluid simulation in this series used the Particle Strength Exchange (PSE) technique to exchange momentum between nearby particles. This requires that the simulation keep track of which particles are near what others -- effectively knowing their nearest neighbors. I presented a simplisitic approach using a uniform grid spatial partition, but others could work -- and this is one of the avenues I encourage people to explore further.

The **stretch and tilt** terms of the vorticity equation describe how vortices interact with each other over distance due to configuration. This is strictly a 3D effect, and it leads to turbulent motion. Without this effect, fluids would behave in a much less interesting way. The algorithms I presented compute this using finite differences but others could work. At the end of this article I mention an interesting side-effect of this computation that could be used to model surface tension.

## Conservation of Mass

The **continuity equation** states that the change in mass of a volume equals inflow/outflow of mass through volume surfaces. As described above, the simulation technique in this series dodged solving that explicitly by imposing that the fluid is incompressible.

## Equation of State

The **equation of state** describes how a fluid expands and contracts (and therefore changes density) due to heat. Coupled with the buoyancy term in the momentum equation, this permitted the algorithm to simulate the intuitive behavior that "hot air rises and cold air sinks".

## Combustion

The **Arrhenius equation** describes how components of fluid transform: fuel to plasma to exhaust. It also describes how fluid heats up. This feeds into the equation of state to model how the fluid density changes with temperature, and hence, causing hot air to rise.

## Drag

Drag describes how fluids interact with solid objects. I presented an approach that builds upon the Particle Strength Exchange (PSE) approach used to model viscosity; I treat fluid particles and solid objects in a similar paradigm. I extended the process to exchange heat too, so solid objects can heat or cool fluids and vice versa.

## *Spatial Discretization*

Fluid equations operate on a continuous temporal and spatial domain, but simulating them on a computer requires discretizing the equations in both time and space. You can discretize space into regions, and those regions could either move (e.g. with **particles**) or not (e.g. with a fixed **grid**).

As its name suggests, the vortex particle method (VPM) is a particle-based method, in contrast to a grid-based method. But the algorithm I presented also uses a uniform grid spatial partition to help answer queries about spatial relationships such as knowing which particles are near which others, or which particles are near solid objects interacting with the fluid. Many spatial partitions are available and within each various implementations are possible. For this article series I chose something reasonably simple and reasonably fast but I suspect that could be improved dramatically, so I describe some ideas to try at the end of this article.

(Other discretizations are possible, for example in a spectral domain. I mention this in passing so curious readers know about other possibilities, but for the sake of brevity I omit details.)

## Vortex Particle Method

In this series, I predominantly employed the Vortex Particle Method for modeling fluid motion. But even within that method you have many specific choices about how to implement various aspects of the numerical solver. Ultimately the computer needs to obtain velocity from vorticity and there are two mathematical approaches to doing so: integral or differential. And each of those mathematical approaches can be solved through multiple numerical algorithms.

The integral techniques I presented include **direct summation** and **treecode**. Direct summation has asymptotic time complexity  $O(N^2)$ , which is the slowest of those presented, but it is also the simplest to implement. Treecode has asymptotic time complexity  $O(N \log N)$ , which is between the slowest and fastest, and has substantially more code complexity than direct summation, but that complexity is worth the speed advantage. Besides those techniques, other options are possible but I did not cover them: For example **multipole methods** have asymptotically low computational complexity order, but mathematically and numerically require much greater complexity.

The differential technique I presented entails solving a vector Poisson equation. Among the techniques I presented, this has the fastest asymptotic runtime, and the math and code are not very complex. Based on that description it seems like the obvious choice, but there is a catch -- involving boundary conditions.

Solving any partial differential equation entails posing boundary conditions -- solving the equations at the spatial bounds of the domain. For integral techniques, the simplest conditions are "open" which is tantamount to having an infinite domain without walls. The simulation algorithm handles solid objects, including walls and floors, which should suffice to impose boundary conditions appropriate for whatever scene geometry interacts with the fluid, so imposing additional boundary conditions would be redundant.

The Poisson solver I presented uses a rectangular box with a uniform grid. It's relatively easy to impose "no-slip" or "no-through" boundary conditions on the box, but then the fluid would move as though it were inside a box. You could move the domain boundaries far from the interesting part of the fluid motion but since the box has a uniform grid, most of the grid cells would have nothing interesting in them -- yet they would cost both memory and compute cycles. So ideally you'd have a Poisson solver that supports open boundary conditions. This is tantamount to knowing the solution at the boundaries -- but the Poisson solver is meant to obtain the solution -- so this is a cyclic dependency.

To solve this problem, I used the integral technique to compute a solution at the domain boundaries (a 2-dimensional surface), then used the Poisson solver to compute a solution throughout the domain interior. This hybrid approach runs in  $O(N)$  time (faster than treecode), and looks better than the treecode solver results.

## Assessment

The Vortex Particle Method works well for fire and smoke. VPM does not work well for liquid-gas boundaries. "Ignored potential flow" cases.

Smoothed Particle Hydrodynamics works well for liquid-gas boundaries but looks viscous.

My attempt to merge them didn't work very well but I still suspect the approach has merit.

## Further Possibilities

The techniques and code I presented in this series provide examples and a starting point for a fluid simulation for video games. To turn this into viable production code would require further refinements to both the simulation and rendering code.

### Simulation

#### Improvements to VPM

I implemented a simplistic uniform grid spatial partitioning scheme. A lot of time is spent performing queries on that data structure. It could be optimized or replaced, for example, by a spatial hash. Also the per-cell container could be switched to a much more lightweight container.

While difficult, it is possible to model those liquid-gas boundaries with a vortex particle method. You could track surfaces using level-sets, use surface geometry to compute curvature, use curvature to compute surface tension, and incorporate those effects into the vorticity equation. Computing curvature entails computing the Hessian, which is related to Jacobian which is already used to compute strain & stress.

The vortex particle method has a glaring mathematical problem. It starts with a bunch of small particles that carry vorticity in a very small region -- so small it is tempting to think of them as points. Vorticity mathematically resembles a magnetic field, and one could draw an analogy between these vortex particles and little tiny magnets. But these magnets would have only a single "pole" -- which is both mathematically and physically impossible. Likewise, there is no such thing as a vortex "point". If you had only a single vortex point, it would be possible for vortex field to have divergence, and this is not mathematically possible nor physically meaningful. And yet, this simulation technique has exactly this problem. One way to solve this problem would be to use vortex filaments, for example topologically forming loops; the vortex loops, being closed, would have no net divergence. (See, for example, "Simulation of Smoke Based on Vortex Filament Primitives" by Angelidis & Neyret .) The filaments could also terminate at fluid boundaries such as at the interfaces with solid objects. The most obvious example of that would be a rotating body -- effectively the body has a vorticity so vortex lines should pass through the body. (Beware that the article "Filament-based smoke with vortex shedding and variational reconnection" as presented at SIGGRAPH 2010 got that wrong -- they had rotating bodies within a fluid but their vortex filaments did not pass through those bodies. They seem to have corrected that in subsequent publications and the YouTube videos that had the error are no longer visible.)

#### Other techniques

Since Smoothed Particle Hydrodynamics is also a fluid simulation technique that uses particles, my intuition is that it should complement the Vortex Particle Method such that some hybrid could work for both wispy and

liquid or goopy fluids. I would not call my attempt successful but I hope it inspires future ideas to unify those approaches -- even though my implementation failed I suspect the basic idea could still be made to work.

This article series did not cover them, but grid-based methods work well in specialized cases, such as where potential flow is important, and for shallow-water waves. And spectral methods are capable of tremendous accuracy, but that is exactly what video games can forsake.

## Rendering

In the code that accompanies these articles, most processing time goes toward rendering, not simulation. That is good news because the simplistic rendering in the sample code does not exploit modern graphics processing units so there is plenty of opportunity to speed that up.

The sample code performs several per-vertex operations, such as computing camera-facing quadrilaterals. That code is embarrassingly parallel, so a programmable vertex shader could execute that on the GPU very quickly because the GPU has hundreds or thousands of processing units.

It turns out, though, that adding more CPU cores to those routines that operate on each vertex does not yield a linear speed-up. This suggests that memory bandwidth limits processing speed; effectively, to speed up processing, the machine would need to access less memory. Again, a solution is readily available: Inside the vertex buffer, instead of storing an element per triangle vertex, store only a single element per particle. It could even be possible to transmit a copy of the particle buffer as-is. Since you can control how the vertex shader accesses memory, that vertex buffer can be in any format you like, including the one the particle buffer has. This implies using less memory bandwidth.

Note that the GPU would likely still need a separate copy of the particle buffer, even if its contents were identical to the particle buffer used by the CPU. The reason is that those processors run asynchronously, so if they shared a buffer, it would be possible for the CPU to modify a particle buffer in the middle of the GPU accessing that data, which could result in inconsistent rendering artifacts. In that case, it might be prudent to duplicate the particle buffer. (But perhaps a DMA engine could make that copy, leaving the CPU unencumbered by that operation.) On the other hand, the visual artifacts of rendering the shared particle buffer might be so small and infrequent that the user might not notice. It's worth trying several variations to find a good compromise between speed and visual quality.

In order for the fluid simulation to look like a continuous and dense fluid instead of a sparse collection of dots, the sample code uses a lot of tracer particles -- tens of thousands. And arguably it would look even better if it had millions of particles. But processing and rendering that is computationally expensive -- both in time and memory. If you used fewer particles of the same size, the rendering would leave gaps. If you increased the particle size, the gaps would close but the fluid would look less wispy -- that is, unless the particles grew only along the direction that smaller particles would appear. There are at least 3 ways to approach this problem: (1) Use volumetric rendering instead of particle rendering. This would involve computing volumetric textures and rendering them with fewer, larger camera-facing quads that access the volumetric texture, and the results can look amazing. (2) Elongate tracer particles in the direction they stretch. One way to do that might be to consider tracers as pairs, where they are initialized very near each other, and are rendered as two ends of a

capsule, instead of treating every particle as an individual blob. You could even couple this with a shader that tracks the previous and current camera transform and introduce a simplistic but effective motion blur; the mathematics is similar for both. (3) Expanding on the idea in option (2), use even more tracers connected together in streaks. For example you could emit tracer particles in sets of 4 (or some  $N$  of your choice) and render those as a ribbon. Note, however, that rendering ribbons can be tricky if the particle cluster "kinks"; it can lead to segments of the ribbon folding such that it has zero area in screen space.

## About the Author

Dr. Michael J. Gourlay works at Microsoft as a principal development lead on HoloLens in the Environment Understanding group. He led the teams that implemented tracking, surface reconstruction and calibration.

He previously worked at Electronic Arts (EA Sports) as the software architect for the Football Sports Business Unit, as a senior lead engineer on *Madden NFL* and original architect of *FranTk* (the engine behind Connected Careers mode), on character physics and *ANT* (the procedural animation system used by EA), on *Mixed Martial Arts*, and as a lead programmer on *NASCAR*. He wrote *Lynx* (the visual effects system used in EA games worldwide) and patented algorithms for interactive, high-bandwidth online applications.

He also developed curricula for and taught at the University of Central Florida, Florida Interactive Entertainment Academy, a top-rated interdisciplinary graduate program that teaches programmers, producers, and artists how to make video games and training simulations.

Prior to joining EA, he performed scientific research using computational fluid dynamics and the world's largest massively parallel supercomputers. Michael received his degrees in physics and philosophy from Georgia Tech and the University of Colorado at Boulder.